

# JDBC, what is it good for?

---

Thomas Risberg  
@trisberg

# About Me

Thomas Risberg (@trisberg)

- Member of the Spring engineering team at Pivotal
- Contributing to Project riff and Spring Cloud Data Flow
- Joined the Spring Framework open source project in 2003 working on JDBC support





# Where it all started ... back in 2002

```
41  * <p>The motivation and design of this class is discussed
42  * in detail in
43  * <a href="http://www.amazon.com/exec/obidos/tg/detail/-/1861007841/">Expert One-On-One J2EE Design and Development</a>
44  * by Rod Johnson (Wrox, 2002).
45  * <br>All SQL issued by this class is logged.
46  * <br>Because this class is parameterizable by the callback interfaces and the
47  * SQLExceptionTranslator interface, it isn't necessary to subclass it.
48  * @author Rod Johnson
49  * @see com.interface21.dao
50  * @version $Id: JdbcTemplate.java,v 1.1 2003/02/11 08:10:22 johnsonr Exp $
51  * @since May 3, 2001
52  */
53  public class JdbcTemplate {
54
55      //-----
56      // Instance data
57      //-----
58      /**
59       * Create a Java 1.4-style logging category.
60       */
61      protected final Logger logger = Logger.getLogger(getClass().getName());
62
63      /**
64       * Used to obtain connections throughout
65       * the lifecycle of this object. This enables this class to
66       * close connections if necessary.
67       */
68      private DataSource dataSource;
69
70      /**
71       * If this variable is false, we will throw exceptions on SQL warnings
72       */
73      private boolean ignoreWarnings = true;
74
75      /** Helper to translate SQL exceptions to DataAccessExceptions */
76      private SQLExceptionTranslator exceptionTranslator;
77  }
```





# Spring solved JDBC try-catch nightmare and more

```
public class JdbcDao {
    private Logger logger = Logger.getLogger("SpringOne");

    public int getBeerCount() {
        Connection conn = null;
        Statement stmt = null;
        ResultSet rs = null;
        int count = 0;
        Properties properties = new Properties();
        try {
            properties.load(new FileInputStream("jdbc.properties"));
        } catch (IOException e) {
            throw new MyDataAccessException("I/O Error", e);
        }
        try {
            Class.forName(properties.getProperty("driverClassName"));
            conn = DriverManager.getConnection(properties.getProperty("url"), properties);
            stmt = conn.createStatement();
            rs = stmt.executeQuery("select count(*) from beers");
            if (rs.next()) {
                count = rs.getInt(1);
            }
        } catch (ClassNotFoundException e) {
            throw new MyDataAccessException("JDBC Error", e);
        } catch (SQLException se) {
            throw new MyDataAccessException("JDBC Error", se);
        } finally {
            if (rs != null) {
                try {
                    rs.close();
                } catch (SQLException ignore) {}
            }
            if (stmt != null) {
                try {
                    stmt.close();
                } catch (SQLException ignore) {}
            }
            if (conn != null) {
                try {
                    conn.close();
                } catch (SQLException ignore) {}
            }
        }
        return count;
    }
}
```

Using  
straight  
JDBC code



## Using Spring

```
public class SpringDao {
    private Logger logger = Logger.getLogger("SpringOne");

    public int getBeerCount() {
        DriverManagerDataSource dataSource = new DriverManagerDataSource();
        int count = 0;
        Properties properties = new Properties();
        try {
            properties.load(new FileInputStream("jdbc.properties"));
        } catch (IOException e) {
            logger.severe(e.toString());
        }
        dataSource.setConnectionProperties(properties);
        dataSource.setDriverClassName(properties.getProperty("driverClassName"));
        dataSource.setUrl(properties.getProperty("url"));

        JdbcTemplate jdbcTemplate = new JdbcTemplate(dataSource);
        count = jdbcTemplate.queryForInt("select count(*) from beers");
        return count;
    }
}
```



From "Spring JDBC" presentation at OSCON 2005

<https://github.com/trisberg/s1p2017-jdbc/tree/master/older-presentations>

# Spring JDBC - who does what?

	Spring	You
DataSource/Connection Configuration		✓
Connection Management	✓	
Specify the SQL statement and its parameters		✓
Manage statement execution and loop through results	✓	
Retrieve and handle data for each row		✓
Manage the transaction	✓	
Translate exceptions to RuntimeExceptions	✓	

# Spring JDBC Support

Choosing an approach for JDBC database access

- **JdbcTemplate** is the classic Spring JDBC approach and the most popular
- **NamedParameterJdbcTemplate** wraps a **JdbcTemplate** to provide named parameters instead of the traditional JDBC "?" placeholders
- **SimpleJdbcInsert** and **SimpleJdbcCall** utilize database metadata to limit the amount of necessary configuration
- RDBMS Objects including **MappingSqlQuery**, **SqlUpdate** and **StoredProcedure** for creating reusable and thread-safe objects during initialization

# Demo

---

## Some JDBC Examples

<https://github.com/trisberg/slp2017-jdbc/tree/master/jdbc-demo>



# What is next for Spring JDBC?

A new **Spring Data JDBC** project

<https://github.com/spring-projects/spring-data-jdbc>

- Spring Data Repository implementation with JDBC
  - **This is NOT an ORM**
- CRUD operations
- Id generation
- NamingStrategy
- Events
  - BeforeDelete / AfterDelete
  - BeforeSave / AfterSave
  - AfterCreation



# Demo

---

## Spring Data JDBC Demo App

<https://github.com/gregturn/spring-data-jdbc-demo>

# Using JDBC in new application architectures

- Microservices / Cloud Deployments
- Event Sourcing / CQRS
- Reactive, Non-Blocking APIs
- Serverless

# JDBC in the Cloud

Connecting to RDBMS databases in cloud environments

- Cloud Foundry
- Kubernetes



# Cloud Foundry for Spring Boot with JDBC

## Running Spring Boot JDBC Apps

- **Java Buildpack** will auto-reconfigure your app injecting its own `DataSource` configuration unless:
  - You have multiple `DataSource` beans defined
  - You have multiple RDBMS services bound to you app
- **Java Buildpack** auto-reconfiguration will also inject a symbolic link to its jdbc driver if you don't provide one in your app jar
- You can turn auto-reconfiguration off with an environment variable:
  - `JBP_CONFIG_SPRING_AUTO_RECONFIGURATION: '[enabled: false]'`
- You probably still want to set the active profile:
  - `SPRING_PROFILES_ACTIVE: cloud`

# Deploying JDBC Spring Boot App on Cloud Foundry

```
cf create-service p-mysql 512mb mysql
cf push -f cloudfoundry/jdbc-demo_manifest.yml \
  -p target/spring-data-jdbc-demo-0.0.1-SNAPSHOT.jar
```

cloudfoundry/jdbc-demo\_manifest.yml

```
applications:
- name: jdbc-demo
  disk_quota: 512M
  instances: 1
  memory: 512M
  routes:
  - route: jdbc-demo.local.pcfdev.io
  services:
  - mysql
  stack: cflinuxfs2
```

resources/application-cloud.properties

```
spring.datasource.platform=mysql
spring.datasource.initialization-mode=always
spring.datasource.driverClassName=org.mariadb.jdbc.Driver
```

resources/schema-mysql.sql

```
CREATE TABLE IF NOT EXISTS Employee ( id BIGINT AUTO_INCREMENT
PRIMARY KEY, firstname VARCHAR(100), lastname VARCHAR(100), role
VARCHAR(20))
CREATE TABLE IF NOT EXISTS Manager ( manager_id BIGINT
AUTO_INCREMENT PRIMARY KEY, NAME VARCHAR(100))
```

# Deploying JDBC Spring Boot App on Kubernetes

```
kubectl apply -f kubernetes/mysql/  
kubectl apply -f kubernetes/
```

kubernetes/jdbc-demo-config.yaml

```
apiVersion: v1  
kind: ConfigMap  
metadata:  
  name: jdbc-demo  
  labels:  
    app: jdbc-demo  
data:  
  application.yaml: |-  
    spring:  
      datasource:  
        url: jdbc:mysql://${MYSQL_SERVICE_HOST}:${MYSQL_SERVICE_PORT}/mysql  
        username: root  
        password: ${MYSQL_ROOT_PASSWORD}  
        driverClassName: org.mariadb.jdbc.Driver  
        testOnBorrow: true  
        validationQuery: "SELECT 1"
```

resources/application-kubernetes.properties

```
spring.datasource.platform=mysql  
spring.datasource.initialization-mode=always  
spring.datasource.driverClassName=org.mariadb.jdbc.Driver
```

resources/schema-mysql.sql

```
CREATE TABLE IF NOT EXISTS Employee ( id BIGINT AUTO_INCREMENT  
PRIMARY KEY, firstname VARCHAR(100), lastname VARCHAR(100), role  
VARCHAR(20))  
CREATE TABLE IF NOT EXISTS Manager ( manager_id BIGINT AUTO_INCREMENT  
PRIMARY KEY, NAME VARCHAR(100))
```



# Event Sourcing / CQRS

- Separates reads from writes
- Uses different data models for Event Store and Read Storage
- Might use different type of data stores for Event Store and Read Storage
- JDBC could be a good for for one or both of them

# JDBC and Async / Reactive

We are not there yet.

Async JDBC proposal from JavaOne 2017

- Developed by the JDBC Expert Group through the Java Community Process
- The API is available for download from OpenJDK at <http://oracle.com/goto/java-async-db>
- Send feedback to [jdbc-spec-discuss@openjdk.java.net](mailto:jdbc-spec-discuss@openjdk.java.net)

Fake it using thread pools, an async layer and Futures to hide your blocking calls

# Serverless JDBC

Running serverless functions accessing RDBMS data

- Spring Cloud Function provides abstraction layer that can be used on:
  - Project riff ✓
  - AWS Lambda ✓
  - Azure Functions
  - Open Whisk
  - Fn Project
  - ...
- Prefer using JDBC rather than JPA/Hibernate for faster cold start



# The JDBC Function

```
public class JdbcWriter implements Function<Map<String, Object>, String> {

    @Autowired
    private JdbcTemplate jdbcTemplate;

    private SimpleJdbcInsert insert;

    @PostConstruct
    public void init() {
        this.insert = new SimpleJdbcInsert(jdbcTemplate)
            .withTableName("data")
            .usingColumns("name", "description")
            .usingGeneratedKeyColumns("id");
    }

    @Override
    public String apply(Map<String, Object> data) {
        logger.info("Received: " + data);
        Object name = data.get("name");
        Object description = data.get("description");
        logger.info("Inserting into data table: [" + name + ", " + description + "]");
        SqlParameterSource input = new MapSqlParameterSource("name", name).addValue("description", description);
        Number newId = insert.executeAndReturnKey(input);
        logger.info("NewId is: " + newId);
        return "{ \"newId\": " + newId + " }";
    }
}
```

# Spring Cloud Function on AWS Lambda

```
AWSTemplateFormatVersion : '2010-09-09'
Transform: AWS::Serverless-2016-10-31
Description: Spring Cloud Function with JDBC Writer
Parameters:
  DBPwd:
    NoEcho: true
    Description: The database account password
    Type: String
Resources:
  Employee:
    Type: AWS::Serverless::Function
    Properties:
      FunctionName: JdbcWriter
      Handler: org.springframework.cloud.function.adapter.aws.SpringBootStreamHandler
      Runtime: java8
      CodeUri: s3://trisberg-functions/jdbc-writer-0.0.1-SNAPSHOT-uber.jar
      Description: Demo JDBC Writer on AWS
      Timeout: 30
      MemorySize: 1024
      Environment:
        Variables:
          SPRING_DATASOURCE_URL: jdbc:mysql://springone.clsmkyl5na.us-east-1.rds.amazonaws.com:3306/test
          SPRING_DATASOURCE_USERNAME: master
          SPRING_DATASOURCE_PASSWORD: { "Ref" : "DBPwd" }
          SPRING_DATASOURCE_PLATFORM: mysql
      Role: arn:aws:iam::641162161031:role/lambda-execution-role
```

# Spring Cloud Function on Project riff

```
apiVersion: projectriff.io/v1
kind: Topic
metadata:
  name: data
---
apiVersion: projectriff.io/v1
kind: Function
metadata:
  name: jdbc-writer
spec:
  protocol: http
  input: data
  container:
    image: trisberg/jdbc-writer:0.0.1-SNAPSHOT
    env:
      - name: FUNCTION_URI
        value: file:///functions/function.jar?handler=example.JdbcWriter
      - name: SPRING_PROFILES_ACTIVE
        value: kubernetes
      - name: SPRING_DATASOURCE_DRIVER_CLASS_NAME
        value: 'org.mariadb.jdbc.Driver'
      - name: SPRING_DATASOURCE_URL
        value: 'jdbc:mysql://data-mysql:3306/mysql'
      - name: SPRING_DATASOURCE_USERNAME
        value: root
      - name: SPRING_DATASOURCE_PASSWORD
        valueFrom:
          secretKeyRef:
            name: data-mysql
            key: mysql-root-password
```



# Demo

---

## Spring Cloud Function JDBC Writer App

<https://github.com/trisberg/slp2017-jdbc/tree/master/jdbc-writer>

# Learn More. Stay Connected.

---

<https://github.com/spring-projects/spring-framework/tree/master/spring-jdbc>

<https://github.com/spring-projects/spring-data-jdbc>

<https://github.com/spring-cloud/spring-cloud-function>

<https://github.com/projectriff/riff>

**SpringOne Platform** by Pivotal.

 @s1p    #springone